

# Improved Dropout for Shallow and Deep Learning

Zhe Li

Joint work with Prof. Gong and Prof. Yang

The University of Iowa

Wednesday 29<sup>th</sup> March, 2017

- 1 Introduction and Problem Setup
- 2 Improved Dropout for Shallow Learning
- 3 Improved Dropout for Deep Learning
- 4 Experimental Results
- 5 Conclusion

# Outline

- 1 Introduction and Problem Setup
- 2 Improved Dropout for Shallow Learning
- 3 Improved Dropout for Deep Learning
- 4 Experimental Results
- 5 Conclusion

# The success of deep learning

- Image Classification

# The success of deep learning

- Image Classification

# The success of deep learning

- Image Classification



→ Cat



→ Dog



→ Goldfinch

# The success of deep learning

- Image Classification



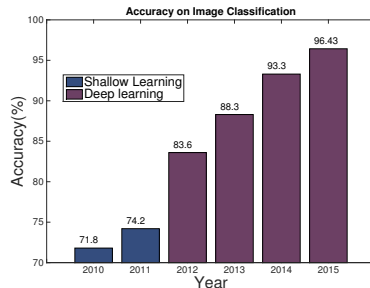
Cat



Dog

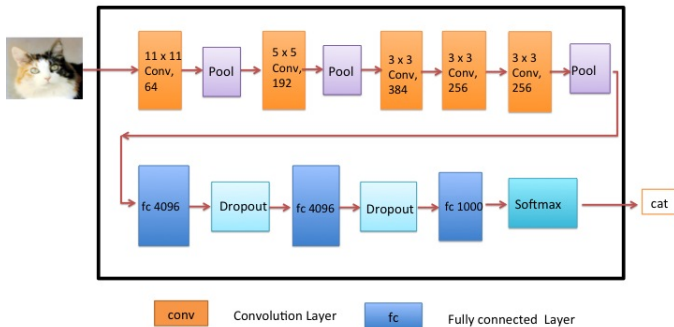


Goldfinch



# Deep Neural Network


The classical example: AlexNet [A Krizhevsky, et .al, 2012]






# Dropout Layer

- Dropout Layer: Uniformly at randomly drop out features.




Before applying dropout




After applying dropout

# Dropout Layer

- Dropout Layer: Uniformly at randomly drop out features.



Before applying dropout




After applying dropout


- Is uniformly dropout optimal?

## Dropout Layer

- Dropout Layer: Uniformly at randomly drop out features.



Before applying dropout




After applying dropout

- Is uniformly dropout optimal?
  - Answered the above question in this work.

# Improved Dropout

## Improved Dropout

- Dropping out the output of the neuron based on multinomial distribution computed from the training data.



**Figure:** Evolutional dropout vs standard dropout on CIFAR100 datasets for deep learning

# Problem Setup

- Let  $(\mathbf{x}, y)$  denote a feature vector and a label, where  $\mathbf{x} \in \mathbb{R}^d$  and  $y \in \mathcal{Y}$ .

# Problem Setup

- Let  $(\mathbf{x}, y)$  denote a feature vector and a label, where  $\mathbf{x} \in \mathbb{R}^d$  and  $y \in \mathcal{Y}$ .
- Denote by  $\mathcal{P}$  the joint distribution of  $(\mathbf{x}, y)$  and by  $\mathcal{D}$  the marginal distribution of  $\mathbf{x}$ .

## Problem Setup

- Let  $(\mathbf{x}, y)$  denote a feature vector and a label, where  $\mathbf{x} \in \mathbb{R}^d$  and  $y \in \mathcal{Y}$ .
- Denote by  $\mathcal{P}$  the joint distribution of  $(\mathbf{x}, y)$  and by  $\mathcal{D}$  the marginal distribution of  $\mathbf{x}$ .
- The goal is to learn a linear prediction function ( $f(\mathbf{x}) = \mathbf{w}^\top \mathbf{x}$ ) that minimizes the expected risk (considering loss function  $\ell(\cdot, y)$ ):

$$\min_{\mathbf{w} \in \mathbb{R}^d} \mathcal{L}(\mathbf{w}) \triangleq \mathbb{E}_{\mathcal{P}}[\ell(\mathbf{w}^\top \mathbf{x}, y)] \quad (1)$$

# Problem Setup

- Denote by  $\epsilon \sim \mathcal{M}$  a dropout noise vector of dimension  $d$ .



# Problem Setup

- Denote by  $\epsilon \sim \mathcal{M}$  a dropout noise vector of dimension  $d$ .
- The corrupted feature vector is given by  $\hat{\mathbf{x}} = \mathbf{x} \circ \epsilon$ , where the operator  $\circ$  represents the element-wise multiplication.

# Problem Setup

- Denote by  $\epsilon \sim \mathcal{M}$  a dropout noise vector of dimension  $d$ .
- The corrupted feature vector is given by  $\hat{\mathbf{x}} = \mathbf{x} \circ \epsilon$ , where the operator  $\circ$  represents the element-wise multiplication.
- Denote by  $\hat{\mathcal{P}}$  the joint distribution of the new data  $(\hat{\mathbf{x}}, y)$  and by  $\hat{\mathcal{D}}$  the marginal distribution of  $\hat{\mathbf{x}}$ .

## Problem Setup

- Denote by  $\epsilon \sim \mathcal{M}$  a dropout noise vector of dimension  $d$ .
- The corrupted feature vector is given by  $\hat{\mathbf{x}} = \mathbf{x} \circ \epsilon$ , where the operator  $\circ$  represents the element-wise multiplication.
- Denote by  $\hat{\mathcal{P}}$  the joint distribution of the new data  $(\hat{\mathbf{x}}, y)$  and by  $\hat{\mathcal{D}}$  the marginal distribution of  $\hat{\mathbf{x}}$ .
- With the corrupted data, the risk minimization becomes

$$\min_{\mathbf{w} \in \mathbb{R}^d} \hat{\mathcal{L}}(\mathbf{w}) \triangleq \mathbb{E}_{\hat{\mathcal{P}}}[\ell(\mathbf{w}^\top (\mathbf{x} \circ \epsilon), y)] \quad (2)$$

# Multinomial Dropout

## Definition 1

A **multinomial dropout** is defined as  $\hat{\mathbf{x}} = \mathbf{x} \circ \epsilon$ , where  $\epsilon_i = \frac{m_i}{kp_i}$ ,  $i \in [d]$  and  $\{m_1, \dots, m_d\}$  follow a multinomial distribution  $Mult(p_1, \dots, p_d; k)$  with  $\sum_{i=1}^d p_i = 1$  and  $p_i \geq 0$ .

# Multinomial Dropout

## Definition 1

A **multinomial dropout** is defined as  $\hat{\mathbf{x}} = \mathbf{x} \circ \epsilon$ , where  $\epsilon_i = \frac{m_i}{kp_i}$ ,  $i \in [d]$  and  $\{m_1, \dots, m_d\}$  follow a multinomial distribution  $Mult(p_1, \dots, p_d; k)$  with  $\sum_{i=1}^d p_i = 1$  and  $p_i \geq 0$ .

- Ability of using non-uniformly sampling probabilities for different features.

# Multinomial Dropout

## Definition 1

A **multinomial dropout** is defined as  $\hat{\mathbf{x}} = \mathbf{x} \circ \epsilon$ , where  $\epsilon_i = \frac{m_i}{kp_i}$ ,  $i \in [d]$  and  $\{m_1, \dots, m_d\}$  follow a multinomial distribution  $Mult(p_1, \dots, p_d; k)$  with  $\sum_{i=1}^d p_i = 1$  and  $p_i \geq 0$ .

- Ability of using non-uniformly sampling probabilities for different features.
- Easy to control the level of dropout by varying the value of  $k$ .

# Multinomial Dropout

- Dropout is a data-dependent regularizer.

# Multinomial Dropout

- Dropout is a data-dependent regularizer.

## Proposition 1

If  $\ell(z, y) = \log(1 + \exp(-yz))$ , then

$$\mathbb{E}_{\hat{\mathcal{P}}}[\ell(\mathbf{w}^\top \hat{\mathbf{x}}, y)] = \mathbb{E}_{\mathcal{P}}[\ell(\mathbf{w}^\top \mathbf{x}, y)] + R_{\mathcal{D}, \mathcal{M}}(\mathbf{w})$$

where  $\mathcal{M}$  denotes the distribution of  $\epsilon$  and

$$R_{\mathcal{D}, \mathcal{M}}(\mathbf{w}) = \mathbb{E}_{\mathcal{D}, \mathcal{M}} \left[ \log \frac{\exp(\mathbf{w}^\top \frac{\mathbf{x}_0 \epsilon}{2}) + \exp(-\mathbf{w}^\top \frac{\mathbf{x}_0 \epsilon}{2})}{\exp(\mathbf{w}^\top \mathbf{x} / 2) + \exp(-\mathbf{w}^\top \mathbf{x} / 2)} \right].$$



# Learning with Multinomial Dropout

## Learning with Multinomial Dropout

# Learning with Multinomial Dropout

## Learning with Multinomial Dropout

- Give the initial solution  $\mathbf{w}_1$ .

# Learning with Multinomial Dropout

## Learning with Multinomial Dropout

- Give the initial solution  $\mathbf{w}_1$ .
- Update the model at  $t^{th}$  iteration:

$$\mathbf{w}_{t+1} = \mathbf{w}_t - \eta_t \nabla \ell(\mathbf{w}_t^\top (\mathbf{x}_t \circ \epsilon_t), y_t) \quad (3)$$

# Learning with Multinomial Dropout

## Learning with Multinomial Dropout

- Give the initial solution  $\mathbf{w}_1$ .
- Update the model at  $t^{th}$  iteration:

$$\mathbf{w}_{t+1} = \mathbf{w}_t - \eta_t \nabla \ell(\mathbf{w}_t^\top (\mathbf{x}_t \circ \epsilon_t), y_t) \quad (3)$$

- Output the final solution:

$$\hat{\mathbf{w}}_n = \frac{1}{n} \sum_{t=1}^n \mathbf{w}_t$$

# Outline

- 1 Introduction and Problem Setup
- 2 Improved Dropout for Shallow Learning**
- 3 Improved Dropout for Deep Learning
- 4 Experimental Results
- 5 Conclusion

## Improved dropout for Shallow Learning

### Theorem 1:

Let  $\mathcal{L}(\mathbf{w})$  be the expected risk of  $\mathbf{w}$  defined in (1). Assume  $\mathbb{E}_{\widehat{\mathcal{D}}}[\|\mathbf{x} \circ \epsilon\|_2^2] \leq B^2$  and  $\ell(z, y)$  is convex and  $G$ -Lipschitz continuous. For any  $\|\mathbf{w}_*\|_2 \leq r$ , by appropriately choosing  $\eta$ , we can have

$$\mathbb{E}[\mathcal{L}(\widehat{\mathbf{w}}_n) + R_{\mathcal{D}, \mathcal{M}}(\widehat{\mathbf{w}}_n)] \leq \mathcal{L}(\mathbf{w}_*) + R_{\mathcal{D}, \mathcal{M}}(\mathbf{w}_*) + \frac{GBr}{\sqrt{n}}$$

How to prove the above theorem?

## Improved dropout for Shallow Learning

### Theorem 1:

Let  $\mathcal{L}(\mathbf{w})$  be the expected risk of  $\mathbf{w}$  defined in (1). Assume  $\mathbb{E}_{\widehat{\mathcal{D}}}[\|\mathbf{x} \circ \epsilon\|_2^2] \leq B^2$  and  $\ell(z, y)$  is convex and  $G$ -Lipschitz continuous. For any  $\|\mathbf{w}_*\|_2 \leq r$ , by appropriately choosing  $\eta$ , we can have

$$\mathbb{E}[\mathcal{L}(\widehat{\mathbf{w}}_n) + R_{\mathcal{D}, \mathcal{M}}(\widehat{\mathbf{w}}_n)] \leq \mathcal{L}(\mathbf{w}_*) + R_{\mathcal{D}, \mathcal{M}}(\mathbf{w}_*) + \frac{GBr}{\sqrt{n}}$$

How to prove the above theorem?

- Standard SGD analysis.

## Improved dropout for Shallow Learning

### Theorem 1:

Let  $\mathcal{L}(\mathbf{w})$  be the expected risk of  $\mathbf{w}$  defined in (1). Assume  $\mathbb{E}_{\widehat{\mathcal{D}}}[\|\mathbf{x} \circ \epsilon\|_2^2] \leq B^2$  and  $\ell(z, y)$  is convex and  $G$ -Lipschitz continuous. For any  $\|\mathbf{w}_*\|_2 \leq r$ , by appropriately choosing  $\eta$ , we can have

$$\mathbb{E}[\mathcal{L}(\widehat{\mathbf{w}}_n) + R_{\mathcal{D}, \mathcal{M}}(\widehat{\mathbf{w}}_n)] \leq \mathcal{L}(\mathbf{w}_*) + R_{\mathcal{D}, \mathcal{M}}(\mathbf{w}_*) + \frac{GBr}{\sqrt{n}}$$

How to prove the above theorem?

- Standard SGD analysis.
- Dropout is a data-dependent regularizer.



## Improved dropout for Shallow Learning

- Minimizing the term  $E_{\hat{\mathcal{D}}}[\|\mathbf{x} \circ \epsilon\|_2^2]$  and the relaxed upper bound of term  $R_{\mathcal{D}, \mathcal{M}}(\mathbf{w}_*)$  yields the optimal sampling probabilities:

$$p_i^* = \frac{\sqrt{E_{\mathcal{D}}[x_i^2]}}{\sum_{j=1}^d \sqrt{E_{\mathcal{D}}[x_j^2]}}, i = 1, \dots, d \quad (4)$$

## Improved dropout for Shallow Learning

- Minimizing the term  $E_{\hat{\mathcal{D}}}[\|\mathbf{x} \circ \epsilon\|_2^2]$  and the relaxed upper bound of term  $R_{\mathcal{D}, \mathcal{M}}(\mathbf{w}_*)$  yields the optimal sampling probabilities:

$$p_i^* = \frac{\sqrt{E_{\mathcal{D}}[x_i^2]}}{\sum_{j=1}^d \sqrt{E_{\mathcal{D}}[x_j^2]}}, i = 1, \dots, d \quad (4)$$

- Can we compute the above probability for dropout?

## Improved dropout for Shallow Learning

- Minimizing the term  $E_{\hat{\mathcal{D}}}[\|\mathbf{x} \circ \epsilon\|_2^2]$  and the relaxed upper bound of term  $R_{\mathcal{D}, \mathcal{M}}(\mathbf{w}_*)$  yields the optimal sampling probabilities:

$$p_i^* = \frac{\sqrt{E_{\mathcal{D}}[x_i^2]}}{\sum_{j=1}^d \sqrt{E_{\mathcal{D}}[x_j^2]}}, i = 1, \dots, d \quad (4)$$

- Can we compute the above probability for dropout?
  - $\times$

## Improved dropout for Shallow Learning

- Practically, we use the empirical second-order statistics to compute the probabilities:

$$p_i = \frac{\sqrt{\frac{1}{n} \sum_{j=1}^n [(\mathbf{x}_j)_i]^2}}{\sum_{i'=1}^d \sqrt{\frac{1}{n} \sum_{j=1}^n [(\mathbf{x}_j)_{i'}]^2}}, i = 1, \dots, d \quad (5)$$

# Outline

- 1 Introduction and Problem Setup
- 2 Improved Dropout for Shallow Learning
- 3 Improved Dropout for Deep Learning**
- 4 Experimental Results
- 5 Conclusion

# Improved dropout for deep learning

- Could we directly use the above idea to Deep Learning?

# Improved dropout for deep learning

- Could we directly use the above idea to Deep Learning?
  - $\times$

# Improved dropout for deep learning

- Could we directly use the above idea to Deep Learning?
  - $\times$
- Why not?



# Improved dropout for deep learning

- Could we directly use the above idea to Deep Learning?
  - $\times$
- Why not?
  - Too expensive to compute dropout probability from all examples.

# Improved dropout for deep learning

- Could we directly use the above idea to Deep Learning?
  - $\times$
- Why not?
  - Too expensive to compute dropout probability from all examples.
- How to address this issue?

# Improved dropout for deep learning

- Could we directly use the above idea to Deep Learning?
  - $\times$
- Why not?
  - Too expensive to compute dropout probability from all examples.
- How to address this issue?
  - Use a mini-batch of examples to calculate the dropout probability.

## Improved dropout for deep learning

- Let  $X^l = (\mathbf{x}_1^l, \dots, \mathbf{x}_m^l)$  denote the outputs of the  $l^{th}$  layer for a mini-batch of  $m$  examples, calculate the probabilities for dropout by

$$p_i^l = \frac{\sqrt{\frac{1}{m} \sum_{j=1}^m [(\mathbf{x}_j^l)_i^2]}}{\sum_{i'=1}^d \sqrt{\frac{1}{m} \sum_{j=1}^m [(\mathbf{x}_j^l)_{i'}^2]}}, i = 1, \dots, d \quad (6)$$

## Improved dropout for deep learning

### Evolutional Dropout for Deep Learning

**Input:** a batch of outputs of a layer:  $X^l = (\mathbf{x}_1^l, \dots, \mathbf{x}_m^l)$   
 and dropout level parameter  $k \in [0, d]$

**Output:**  $\hat{X}^l = X^l \circ \Sigma^l$

Compute sampling probabilities by (6)

For  $j = 1, \dots, m$

Sample  $\mathbf{m}_j^l \sim \text{Mult}(p_1^l, \dots, p_d^l; k)$

Construct  $\epsilon_j^l = \frac{\mathbf{m}_j^l}{k\mathbf{p}^l} \in \mathbb{R}^d$ , where  $\mathbf{p}^l = (p_1^l, \dots, p_d^l)^\top$

Let  $\Sigma^l = (\epsilon_1^l, \dots, \epsilon_m^l)$  and compute  $\hat{X}^l = X^l \circ \Sigma^l$


**Figure:** Evolutional Dropout applied to a layer over a mini-batch

# Outline

- 1 Introduction and Problem Setup
- 2 Improved Dropout for Shallow Learning
- 3 Improved Dropout for Deep Learning
- 4 Experimental Results**
- 5 Conclusion

## Experimental Results for Shallow Learning

Training/test error between standard and improved dropout



**Figure:** data-dependent dropout vs. standard dropout on three datasets (real-sim, news20 and RCV1) for logistic regression

# Experimental Results for Deep Learning

- Implemented in CudaConvNet Library.



# Experimental Results for Deep Learning

- Implemented in CudaConvNet Library.
- Using four benchmark datasets: MNIST, SVHN, CIFAR10, CIFAR100.


# Experimental Results for Deep Learning

- Implemented in CudaConvNet Library.
- Using four benchmark datasets: MNIST, SVHN, CIFAR10, CIFAR100.
- Different neural network structures from the existing literatures.

# Experimental Results for Deep Learning

- Implemented in CudaConvNet Library.
- Using four benchmark datasets: MNIST, SVHN, CIFAR10, CIFAR100.
- Different neural network structures from the existing literatures.
- Training strategy.

## Experimental Results for Deep Learning



**Figure:** Evolutional dropout vs. standard dropout on four benchmark datasets (MNIST, SVHN, CIFAR-10 and CIFAR-100) for deep learning

## Experimental Results for Deep Learning

Compared to Batch Normalization




Figure: Evolutional dropout vs BN on CIFAR-10.

# Outline

- 1 Introduction and Problem Setup
- 2 Improved Dropout for Shallow Learning
- 3 Improved Dropout for Deep Learning
- 4 Experimental Results
- 5 Conclusion**

# Conclusion

- Proposed a multinomial dropout for shallow learning.

# Conclusion

- Proposed a multinomial dropout for shallow learning.
- Demonstrated that this proposed distribution-dependent dropout leads to a faster convergence and a smaller generalization error through the risk bound analysis.



# Conclusion

- Proposed a multinomial dropout for shallow learning.
- Demonstrated that this proposed distribution-dependent dropout leads to a faster convergence and a smaller generalization error through the risk bound analysis.
- Proposed an efficient evolutionary dropout for deep learning.

# Conclusion

- Proposed a multinomial dropout for shallow learning.
- Demonstrated that this proposed distribution-dependent dropout leads to a faster convergence and a smaller generalization error through the risk bound analysis.
- Proposed an efficient evolutionary dropout for deep learning.
- Justified the proposed dropouts for both shallow and deep learning empirically.

# Question?